

Welcome to DATA1030: Hands-on Data Science!

Instructor: Andras Zsom

HTA: Pooja Barai

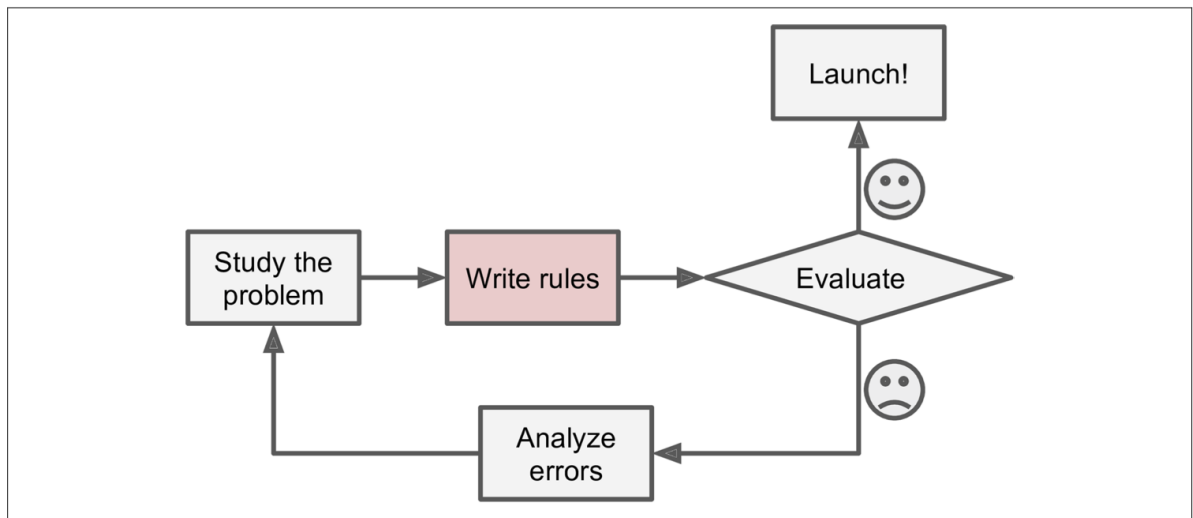
TAs: Tongtong Zhao, Xufan Wang, Yingfei Hong,
Yuetian Li

The goal of this course: supervised Machine Learning (ML)

- supervised ML is probably the most successful area in ML (based on value created)
 - **online advertising**: given an ad and user info, will the user click on the ad?
 - **real estate**: given home features, can we predict the house price?
 - **finance**: given an applicant and a financial product (e.g., a loan), will this applicant be able to successfully pay back the loan?
 - **health care**: given a patient, symptoms, and maybe test results, can we predict the illness?
 - ...
- supervised ML pros:
 - **automation**: computers perform calculations faster than humans (and computers are cheaper)
 - **learn from examples**: no need to explicitly tell the computer what to do. the computer figures out what to do based on examples (data)
- supervised ML con:
 - it can be difficult or labor-intensive to collect training data
 - there is no guarantee that you will be able to develop an accurate model based on the data you have

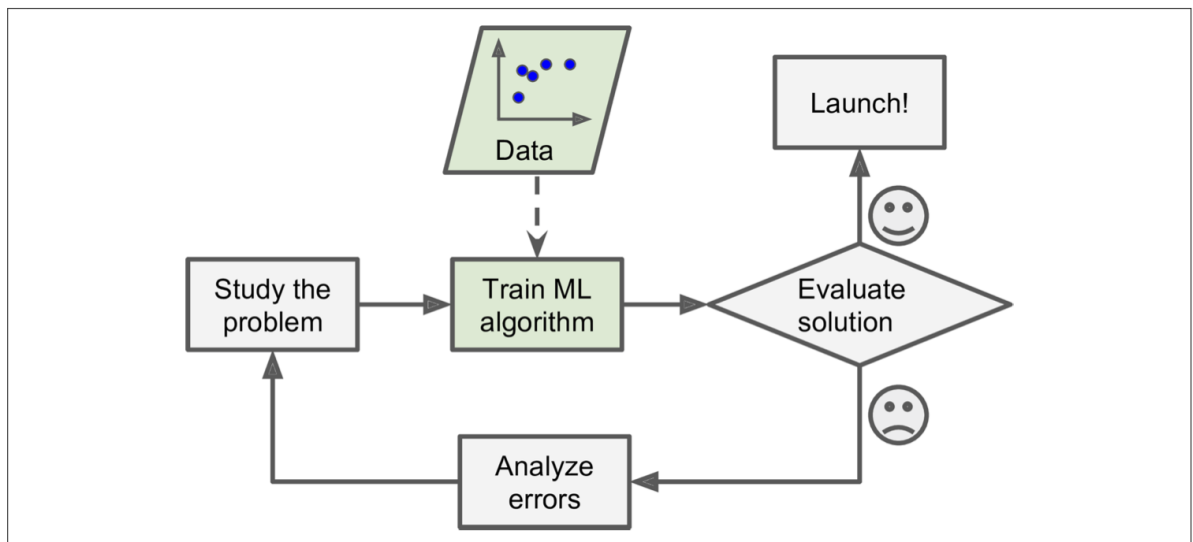
Example: spam filters

- Traditional coding pipeline with explicit instructions



Example: spam filters

- ML pipeline




- the data: feature matrix (X) and target variable (Y)
 - X can be structured (tabular data most commonly stored in excel and csv files or SQL databases)
 - X can be unstructured (e.g., images, text, voice recording, video)
 - Y can be categorical, the problem is **classification** (e.g., click or not click on an ad, sick or not sick)
 - Y can be continuous, the problem is **regression** (e.g., predict house price, stock price, age)
- we focus on structured data during this class!

Structured data

|data_point_n|x_n1|x_n2|...|x_nj|...|x_nm|y_n|

Other areas of ML

- unsupervised ML
 - only the feature matrix X is available, there is no target variable
 - the goal is to find structure (clusters) in the data
 - often used in customer segmentation
- recommender systems
 - recommend products to a customer based on what products similar customers enjoyed
- reinforcement learning
 - the learning system, called an agent, can observe the environment, select and perform actions, and get rewards and penalties in return. Goal: come up with strategy to maximize rewards
 - often used when virtual environment is available (e.g., games like go or warcraft)
 - sounds appealing to use in real environments (like self-driving cars) but agents learn slow, lots of cars would need to be broken to teach an agent to drive this way
- deep learning
 - uses neural networks and often works with unstructured data
 - technically deep learning is supervised or unsupervised
 - extremely successful on large datasets



Quiz

Learning objectives

By the end of the semester, you will be able to

- explore and visualize the dataset,
- develop a ML pipeline from scratch to deployment,
- make data-driven decisions during the pipeline development,
- handle non-standard ML problems like missing data, non-iid data,
- provide explanations with your model,
- explain your findings to technical and non-technical audiences.

A few words about python

- widely used in data science because of sklearn, pandas, deep learning packages
 - packages are easy to (mis)use
- relatively easy to write code but difficult to write computationally efficient code
 - the divide between package developers and users is huge!
 - you will need to spend a lot of time reading the manuals and verifying results
- the lecture notes contain code that has been tested
 - this is misleading!
 - I spent a lot of time testing the code but I deleted those lines to keep the final code clean
 - but when you write code, you should absolutely PRINT ALL VARIABLES and TEST EVERY SINGLE LINE!
 - you will learn how to interpret error messages and how to debug your code
- test-driven code development is encouraged
 - first come up with a test
 - create a couple of test cases with known results
 - i.e., if my code does what I think it should, I'll get a certain output given certain input
 - then write the code

Course structure

Canvas: <https://canvas.brown.edu/courses/1089420>

Course components:

- lectures
 - in person but recordings will be posted on canvas
- weekly problem sets

- coding problems and questions with 1-2 paragraph answers
 - the questions prepare you for your job interviews
- one final exam
 - most likely short, open-book exam on December 14th, 2pm
- one semester-long project
 - find a dataset and come up with your own machine learning question
 - develop code individually, but feel free to discuss with others
 - assigned TA mentor with regular dedicated meetings

Grading

- lectures: **10%** weight
 - you can skip lectures three times during the term
 - the quiz answers are not graded so don't worry if you don't answer them correctly!
- weekly problem sets: **35%** weight
 - this is a substantial component
- exam: **20%** weight
- project: **35%** weight
 - important component!
 - make sure to spend sufficient time on this each week!
 - the semester will go by very quickly...
- **90% minimum is necessary to get an A** but I reserve the right to lower the threshold
- my experience is that Bs are rare, C is given under exceptional circumstances

Project

- look for datasets on the [UCI Machine Learning Repository](#), on [Kaggle](#), or google's [dataset search engine](#).
- Bring your own dataset!
 - if you have your own dataset you'd like to work with, this is the perfect opportunity!
- Avoid the most popular datasets!
 - no Titanic, no iris for example
- avoid these four datasets because we will use them in class and you'll work with them in the problem sets
 - [adult dataset](#)
 - [kaggle house price dataset](#)
 - [hand postures dataset](#)
 - [diabetes dataset](#)
- work on a classification or regression problem!
- start looking for datasets now and talk to the TAs or come to my office hours if you have questions!

Rough deadlines

- **1st project presentation:** early/mid October (multiple dates)
 - short presentation on dataset, EDA, and ML question (7 min + 2 min questions per student)
 - rubric will be available two weeks in advance
- **1st project progress report:** early/mid October
 - dataset selection, EDA, and formulate your ML question
 - rubric will be available two weeks in advance
- **final presentations:** early December
 - another short presentation on ML pipeline and results
 - rubric will be available two weeks in advance
- **final project report:** early December
 - the complete ML pipeline and results
 - rubric will be available two weeks in advance
- **final exam:** December 14th, 2pm
- grades finalized and submitted by December 16th

Other course resources

- Ed discussion: course forum
 - feel free to discuss any questions or concerns regarding the material
 - please post publicly whenever possible (but you can still post anonymously)
 - if you have a question, it is likely that multiple students have the same question
 - the TAs and I will keep an eye on it and answer questions in a timely manner
 - disclaimer: I turn off my laptop after 5pm and during the weekends
- office hours (TAs and mine)
 - see the course google calendar
- An Introduction to Statistical Learning ([book](#))
- Introduction to Machine Learning with Python ([book](#))
- Harry Potter and the Methods of Rationality ([fan fiction](#) by Eliezer Yudkowski)
 - half joking, half serious about this one :)

Mud card

In []: